

Advance SQL

Accessing SQL from programming language

- Some complex queries cannot be written completely using SQL commands.
 - Such queries can be expressed in a language such as C, Java, etc.
- Non-declarative actions
 - Printing a report
 - Interacting with a user
 - Sending a results to GUI

Accessing SQL from programming language

- Two Approaches
 - Dynamic SQL
 - Embedded SQL
- Dynamic SQL
 - At runtime
 - SQL queries as character strings, submit to the database, and then retrieve the response
- Embedded SQL
 - SQL queries are identified at the compile time using a pre-processor
 - Pre-processor submits the queries to the database for precompilation and optimization
 - Replace the SQL queries with the optimized code

APIs

- Almost, each database system provides APIs to interact with the database
 - Connect to the database
 - Send SQL commands to the database
 - Fetch tuples of results one-by-one into program
- Tools
 - ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic. Other API's such as ADO.NET sit on top of ODBC
 - JDBC (Java Database Connectivity) works with Java

ODBC

- Open DataBase Connectivity
 - Standard for application programs to communicate with a database
- Applications such as spreadsheets, GUI, etc. can use ODBC

JDBC

- Java DataBase Connectivity
 - JAVA API to interact with the database supporting JDBC
- Supports metadata retrieval
 - Check whether a relation exists or not
 - Names, types of attributes
- JDBC model for communication
 - Open a connection
 - Create a statement object
 - Execute queries using the statement object to send queries and fetch results
 - Exception mechanism to handle errors

JDBC - An example

- Open connection

```
Class.forName ("oracle.jdbc.driver.OracleDriver");           //Loading database driver
```

```
Connection conn = DriverManager.getConnection ("jdbc:oracle:thin@server.port:db",  
                                              userid, passwd);
```

- Create statement

```
Statement stmt = conn.createStatement();
```

JDBC - An example

- Execute statement

- executeUpdate - (insert, update, delete)

```
stmt.executeUpdate("insert into instructor values('77987', 'Kim', 'Physics', 98000)");
```

- Retrieving the Result of a Query

- executeQuery - returns a relation (select)

```
ResultSet rset = stmt.executeQuery(  
    "select dept_name, avg(salary)" + " from instructor" + " group by dept_name");  
  
while (rset.next())  
{  
    System.out.println(rset.getString("dept_name") + ", " + rset.getFloat(2));  
}
```


JDBC - An example

- Closing connection

```
stmt.close();  
conn.close();
```

JDBC - An example

```
public static void JDBCexample(String userid, String passwd) {  
    try {  
        Class.forName ("oracle.jdbc.driver.OracleDriver");  
        Connection conn = DriverManager.getConnection ("jdbc:oracle:thin@server:port:db", userid, passwd);  
        Statement stmt = conn.createStatement();  
        try {  
            stmt.executeUpdate("insert into instructor values('77987', 'Kim', 'Physics', 98000)");  
        } catch ( SQLException sqle) { System.out.println("Could not insert tuple. " + sqle); }  
  
        ResultSet rset = stmt.executeQuery("select dept-name, avg (salary)+"from instructor"+"group by dept-name");  
  
        while (rset.next()) { System.out.println(rset.getString("dept-name") + " " + rset.getFloat(2)); }  
        stmt.close(); conn.close();  
    }  
    catch (Exception sqle) { System.out.println("Exception : " + sqle); }  
}
```

Prepared Statement

- If some values are not known at the moment, but will be available at a later stage
 - Create statement when then values are available, 'OR'
 - Prepare statement and supply the values when available
 - Useful, when same query with different values to be executed multiple times

```
PreparedStatement pStmt = conn.prepareStatement("insert into instructor values(?,?,?,?)");  
//.....  
pStmt.setString(1, "88877");  
pStmt.setString(2, "Perry");  
pStmt.setString(3, "Finance");  
//.....  
pStmt.setInt(4, 125000);  
pStmt.executeUpdate();
```

SQL Injections and Prepared Statement

- Handles special character in input efficiently

“select * from instructor where *name* = ‘ “ + name + ” ’ ”

- Case 1:

- What if user enters O'Henry as name?

- Syntax error

- `setString()` of `preparedStatement` handles this scenario by inserting escape characters (`\'`)

- Case 2:

- What if user enters X' or 'Y' = 'Y as name?

select * from instructor where *name* = 'X' or 'Y' = 'Y'

- `setString()`

select * from instructor where *name* = 'X\' or \'Y\' = \'Y'

Metadata - Relation

- ResultSetMetaData object allows to access the metadata of a relation returned by executeQuery()
- Name and Type of the columns

```
ResultSet rs = executeQuery(".....");
ResultSetMetaData rsmd = rs.getMetaData();
for(int i = 1; i <= rsmd.getColumnCount(); i++)
{
    System.out.println(rsmd.getColumnName(i));
    System.out.println(rsmd.getColumnTypeName(i));
}
```

Metadata - Database

- DatabaseMetaData object allows to access the metadata of the database
- Columns

```
DatabaseMetaData dbmd = conn.getMetaData();  
ResultSet rs = dbmd.getColumns(null, "univdb", "department", "%");  
// getColumns(Catalog, Schema-pattern, Table-pattern, Column-Pattern)  
// returns: "COLUMN_NAME", "TYPE_NAME"
```

Metadata - Database

- **Tables**

```
ResultSet rs = dbmd.getTables ("", "", "%", new String[] {"TABLE"})  
// getTables(Catalog, Schema, Table-name, Table-type (table, view, etc.))  
// returns: "TABLE_NAME", "TABLE_TYPE"
```

- **Primary Keys**

```
ResultSet rs = dbmd.getPrimaryKeys( catalog, schema, tableName);  
// getTables(Catalog, Schema, Table-name, Table-type (table, view, etc.))  
// returns: "COLUMN_NAME", "KEY_SEQ"
```

Transactions in JDBC

- By default, each SQL statement is treated as a separate transaction that is committed automatically
- Turn off automatic commit on a connection
 - `conn.setAutoCommit(false);`
- Transactions must then be committed or rolled back explicitly
 - `conn.commit();` or
 - `conn.rollback();`
- Turn off automatic commit on a connection
 - `conn.setAutoCommit(true)`

Embedded SQL

- A language in which SQL queries are embedded is referred to as a host language

EXEC SQL <embedded SQL statement >;

- In COBOL, the semicolon is replaced with END EXEC
- In Java, embedding uses # SQL { };

Embedded SQL

- Connect to database

EXEC SQL **connect to server user** *user_name* **using** *password*;

- Variables

- Program variables can be used; however, they are preceded by colon (:)
 - *:credit-amount*
- They must be declared before usage

EXEC SQL **BEGIN DECLARE SECTION**;

int *credit-amount*; // Host language syntax;

EXEC SQL **END DECLARE SECTION**;

Embedded SQL

- Relational query

declare c cursor for <SQL query>;

variable c identifies the query

- Cursor: A *temporary area for work* in memory system while the execution of a statement is done.
- Results will be computed on **open** and **fetch** commands

- Query: Find *name* and *id* of the student with credits more than *:credit-amount*

EXEC SQL

declare c cursor for

select *ID, name* **from** *student* **where** *tot_credit > :credit-amount ;*

Embedded SQL

- Open

- Evaluates the query and stores the results in temporary area

EXEC SQL **open** *c*;

- Fetch

- From temporary area to host language variables
- One variables for each attribute
- Fetched one tuple at a time

EXEC SQL **fetch** *c into* *:si*, *:sn*;

:si and *:sn* are the program variables to store *id* and *name*, respectively.

Embedded SQL

- **Close**

- The close statement causes the database system to delete the temporary relation that holds the result of the query.

EXEC SQL **close** c;

- **SQL communication area (SQLCA)**

- If the SQL query results in an error, the database system stores an error diagnostic in the SQLCA variables.
- On fetch, when no further tuples remain to be processed, the character array variable SQLSTATE in the SQLCA is set to '02000' (meaning “no more data”)

Embedded SQL - Update/Delete

- Can update tuples fetched by cursor by declaring that the cursor is for update

EXEC SQL

declare c cursor for

select * from *instructor* where dept_name = 'Music'

for update

- We then iterate through the tuples by performing fetch operations on the cursor (as illustrated earlier), and after fetching each tuple we execute the following code:

EXEC SQL

update *instructor* set salary = salary + 1000

where current of c;

Function and procedure

- Procedures and functions allow “business logic” to be stored in the database and executed from SQL statements.
 - how many courses a student can take in a given semester,
 - the minimum number of courses a full-time instructor must teach in a year,
 - the maximum number of majors a student can be enrolled in,
- Rather than ensuring such business logic in the programming language, define them as the stored procedure in the database.

SQL Function

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (d_name varchar(20))  
returns integer  
begin  
    declare d_count integer;  
    select count(*) into d_count from instructor  
    where instructure.dept_name = d_name  
    return d_count;  
end
```


SQL Function

- The function `dept_count` can be used to find the department names and budget of all departments with more than 12 instructors.

`select dept_name, budget from department where dept_count(dept_name) > 12`

- Compound statement: `begin ... end`
 - May contain multiple SQL statements between `begin` and `end`.
- **returns:** indicates the variable-type that is returned (e.g., integer)
- **return:** specifies the values that are to be returned as result of invoking the function

Table function

- The SQL standard supports functions that can return tables as results

```
create function instructors_of (dept_name varchar(20))  
returns table (  
    ID varchar(5), name varchar(20), dept_name varchar(20), salary numeric (8,2))  
return table (  
    select ID, name, dept_name, salary  
from instructor where instructor.dept_name = instructors_of.dept_name);
```

- Get the instructors' detail from finance department

```
select * from table(instructors_of('Finance'));
```

SQL Procedure

- The dept_count function could instead be written as procedure:

```
create procedure dept_count_proc (in dept_name varchar(20), out d_count integer)
begin
    select count(*) into d_count from instructor
    where instructor.dept_name = dept_count_proc.dept_name
end
```

- Keywords **in** and **out**
 - Parameter expected and parameters to return
- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the call statement.

```
declare d_count integer;
call dept_count_proc( 'Physics', d_count);
```

Language constructs for function and procedure

- SQL also supports other programming language constructs, that gives it almost all the power of general-purpose language
 - while, for, etc.
- Variables can be declared using **declare** statement and assigned using **set**
- Compound statement:
 - May contains multiple SQL commands between **begin** and **end** statements
 - Local variables can also be declared
- Compound statements as transaction
 - **begin atomic end**

Loops - While/Repeat/For

```
while boolean_expression do  
    sequence of statements;  
end while
```

```
repeat  
    sequence of statements;  
until boolean_expression  
end repeat
```

```
declare n integer default 0;  
for r as  
    select budget from department where dept name = 'Music'  
do  
    set n = n - r.budget  
end for
```

Conditional statements: if-then-else

```
if boolean expression  
then  
    statement or compound statement  
elseif boolean expression  
then  
    statement or compound statement  
else  
    statement or compound statement  
end if
```

Registers a student to a course after ensuring classroom capacity is not exceeded

```
create function registerStudent(in s_id varchar(5), ..., out errorMsg varchar(100))  
returns integer  
begin  
    declare currEnrol int;  
        select count(*) into currEnrol from takes where ... // get current count  
    declare limit int;  
        select capacity into limit from classroom ... // get capacity of course  
    if (currEnrol < limit)  
        begin  
            insert into takes values (s_id, ....., );  
            return(0); // success  
        end  
    set errorMsg = 'Enrollment limit reached for course';  
    return(-1); // failure  
end;
```

Other constructs

- **leave** - to exit the loop (break)
- **iterate** - starts the next tuple from the beginning of the loop (continue)
- Exception handling

```
    declare out_of_classroom_seats condition
    declare exit handler for out_of_classroom_seats
begin
    ...
    signal out_of_classroom_seats
    ...
end
```

The handler here is **exit** -- causes enclosing begin..end to be exited

External language routines

- Procedural extensions to SQL have some drawbacks
 - Efficiency
 - Different database have different formats
- SQL permits the use of functions and procedures written in other languages such as C, C++, Java, etc.
 - Can be more efficient than function defined in SQL
 - Computations that cannot be carried out in SQL can be executed by these functions.

External language routines

```
create procedure dept_count_proc(in dept_name varchar(20), out count integer)  
language C  
external name ' /usr/avi/bin/dept_count_proc'
```

```
create function dept_count(dept_name varchar(20))  
returns integer  
language C  
external name '/usr/avi/bin/dept_count'
```

External language routines

- Drawback
 - Functions defined and compiled outside the database system may be loaded and executed with the database-system code.
 - risk of accidental corruption of database structures
 - security risk, allowing users access to unauthorized data
- Direct execution in the database system's space is used when efficiency is more important than security.

External language routines

- Solution to security risk, **sandbox**
 - Use a safe language like Java, C#, etc. which cannot be used to access/damage other parts of the database code.
 - Sandbox allows Java, C# code to access its own memory area, and prevents them to access the memory of query execution process
 - Parameters and results communicated via inter-process communication